

# Augmented Reality on FPGA

## Realtime Object Recognition and Image Processing

Logan P. Williams   José E. Cruz Serrallés

15 November 2011

- Overlay a digital image on a physical object in realtime.
- In this case, we want to identify a picture frame in captured video, and output video with another image distorted to fit on top of the picture frame.

# Example Image



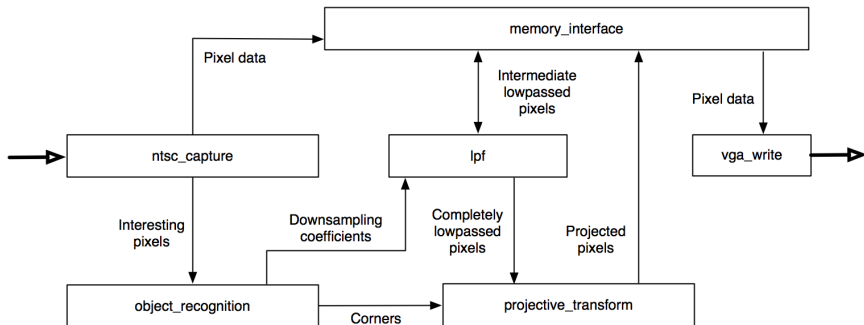
# Example Image




# Example Image



# Top-Level Overview




# projective\_transform: Purpose



- Skew to any arbitrary convex quadrilateral


# projective\_transform: How the algorithm works



- 1 Calculate the distance of line  $\overline{A'D'}$  and assign it to  $d_{ad}$ .
- 2 Do the same for  $\overline{B'C'}$  and assign it to  $d_{bc}$ .
- 3 Create two “iterator points,” point  $I_A$  and  $I_B$  initially located at  $A'$  and  $B'$ .
- 4 Let  $o_x = 0$  and  $o_y = 0$
- 5 Calculate the distance between the iterator points, assign it to  $d_i$ .
- 6 Create a third iterator point,  $I_C$  at the location  $I_A$ .




# projective\_transform: How the algorithm works



- 7 Assign the pixel value of  $I_C$  to pixel  $(o_x, o_y)$  in the original image.
- 8 Move  $I_C$  along line  $\overline{I_A I_B}$  by an amount  $= \frac{d_i}{width_{original}}$ .
- 9 Increment  $o_x$ .
- 10 Repeat steps 7–9 until  $I_C = I_B$ .

# projective\_transform: How the algorithm works




- 11 Move  $I_A$  along line  $\overline{A'D'}$  by an amount  $= \frac{d_{ad}}{height_{original}}$ .
- 12 Move  $I_B$  along line  $\overline{B'C'}$  by an amount  $= \frac{d_{bc}}{height_{original}}$ .
- 13 Increment  $o_y$ .
- 14 Repeat steps 5–13 until  $I_A = D'$  and  $I_B = C'$ .

# projective\_transform: Example




Figure: The original image


# projective\_transform: Example



# projective\_transform: Example



# projective\_transform: Example



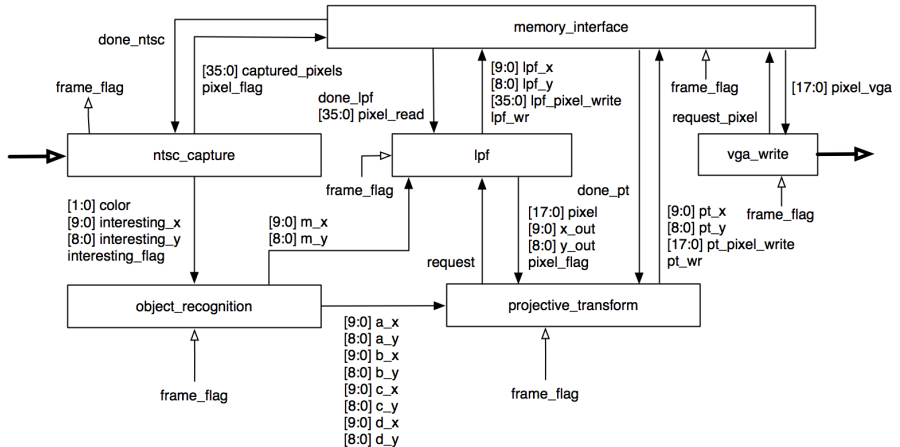
# projective\_transform: Example



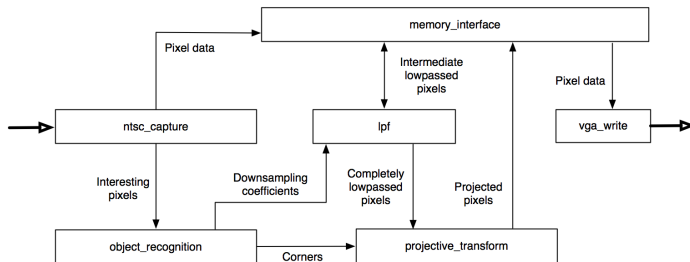
- Straightfoward implementation of the above algorithm
- Uses coregen Divider modules for the divisions
- Requires only  $2 \times 640 \times 480 + 4 \times 480$  multiplications per clock cycle
- Uses an iterative algorithm for finding distances (pipelined at the end of each line of the image)
- Processes pixels “on-the-fly” from  $LPF$
- Negligible memory requirements (a handful of registers)



# projective\_transform: How it Interfaces



# object\_recognition



- Mark corners of frame with four differently colored dots.
- Recognition begins in the `ntsc_capture` module, which detects these colors as it is capturing data and sends the pixel info to the `object_recognition` module.

- Take linear weighted center of mass for each image
- Sums the (x,y) coordinates for each color as it receives them. (8 running sums, 2 for each color)
- When the frame is done, divide each sum by the number of summed items
- The resulting 4 (x,y) pairs are the corners of the frame
- By looking for pixels in `ntsc_capture` we significantly reduce the amount of time spent in `object_recognition`

# LPF: its purpose

- `projective_transform` →  
aliasing

graphic showing normal signal

# LPF: its purpose

- `projective_transform` →  
aliasing

graphic aliases

# LPF: its purpose

- `projective_transform` → aliasing
- Aliasing reduces the quality of an image

zoom in on aliased pixels

# LPF: its purpose

- `projective_transform`  $\rightarrow$  aliasing
- Aliasing reduces the quality of an image
- Lowpass filtering prevents aliasing

picture depicting lowpass filter in 2D

# LPF: its purpose

- `projective_transform` → aliasing
- Aliasing reduces the quality of an image
- Lowpass filtering prevents aliasing
- Information of an image is mostly phase

picture of original picture



# LPF: its purpose

- `projective_transform` → aliasing
- Aliasing reduces the quality of an image
- Lowpass filtering prevents aliasing
- Information of an image is mostly phase

picture of other image

# LPF: its purpose

- `projective_transform` → aliasing
- Aliasing reduces the quality of an image
- Lowpass filtering prevents aliasing
- Information of an image is mostly phase

picture of original phase with  
other's magnitude

# LPF: its purpose

- `projective_transform` → aliasing
- Aliasing reduces the quality of an image
- Lowpass filtering prevents aliasing
- Information of an image is mostly phase
- Symmetric Type I FIR filter → 0 phase distortion

picture of original phase with  
other's magnitude

# LPF: its purpose

- `projective_transform` → aliasing
- Aliasing reduces the quality of an image
- Lowpass filtering prevents aliasing
- Information of an image is mostly phase
- Symmetric Type I FIR filter → 0 phase distortion
- Parks-McClellan: reasonable accuracy, symmetric, easily calculable

frequency response of  
Parks-McClellan filter

# LPF: its purpose

- `projective_transform` → aliasing
- Aliasing reduces the quality of an image
- Lowpass filtering prevents aliasing
- Information of an image is mostly phase
- Symmetric Type I FIR filter → 0 phase distortion
- Parks-McClellan: reasonable accuracy, symmetric, easily calculable
- FIR PM filter reduces mem. accesses to 1.5/pixel

- 1 Given an arbitrary image & skewing coefficients  $M_x$  &  $M_y$ .

graphic showing the interface  
between object\_recognition and  
LPF  
image  
magnitude fourier plot of image

# LPF: the algorithm

- 1 Given an arbitrary image & skewing coefficients  $M_x$  &  $M_y$ .
- 2 Fetch a filter with cutoff  $\frac{\pi}{M_y}$ .

magnitude plot of image  
magnitude plot of filter with cutoff  
pi/2

# LPF: the algorithm

- 1 Given an arbitrary image & skewing coefficients  $M_x$  &  $M_y$ .
- 2 Fetch a filter with cutoff  $\frac{\pi}{M_y}$ .
- 3 Filter each column and store in memory.

magnitude plot of image  
magnitude plot of filter with cutoff  
 $\pi/2$   
magnitude fourier plot of filtered



# LPF: the algorithm

- 1 Given an arbitrary image & skewing coefficients  $M_x$  &  $M_y$ .
- 2 Fetch a filter with cutoff  $\frac{\pi}{M_y}$ .
- 3 Filter each column and store in memory.
- 4 Fetch a filter with cutoff  $\frac{\pi}{M_x}$ .

magnitude plot of filtered image  
magnitude plot of filter with cutoff  
pi/4

# LPF: the algorithm

- 1 Given an arbitrary image & skewing coefficients  $M_x$  &  $M_y$ .
- 2 Fetch a filter with cutoff  $\frac{\pi}{M_y}$ .
- 3 Filter each column and store in memory.
- 4 Fetch a filter with cutoff  $\frac{\pi}{M_x}$ .
- 5 Filter each row and output to `projective_transform`.

magnitude plot of filtered image  
magnitude plot of filter with cutoff  
 $\pi/5$   
magnitude plot of output

# LPF: the algorithm

- 1 Given an arbitrary image & skewing coefficients  $M_x$  &  $M_y$ .
- 2 Fetch a filter with cutoff  $\frac{\pi}{M_y}$ .
- 3 Filter each column and store in memory.
- 4 Fetch a filter with cutoff  $\frac{\pi}{M_x}$ .
- 5 Filter each row and output to `projective_transform`.
- 6 Repeat this process every refresh cycle.

magnitude plot of original  
magnitude plot of filter  
magnitude plot of output

- 1 image is a lot of data:  
 $640 \cdot 480 \cdot 24 \text{ bits} \approx 0.88\text{MiB}$

- 1 image is a lot of data:  
 $640 \cdot 480 \cdot 24 \text{ bits} \approx 0.88\text{MiB}$
- Total BRAM on board: 0.316MiB

- 1 image is a lot of data:  
 $640 \cdot 480 \cdot 24 \text{ bits} \approx 0.88\text{MiB}$
- Total BRAM on board: 0.316MiB
- We need to store 4 images in memory!

- 1 image is a lot of data:  
 $640 \cdot 480 \cdot 24 \text{ bits} \approx 0.88\text{MiB}$
- Total BRAM on board: 0.316MiB
- We need to store 4 images in memory!
- Let's use the ZBT RAM:  $2 \cdot 2.25\text{MiB}$

- 1 image is a lot of data:  
 $640 \cdot 480 \cdot 24 \text{ bits} \approx 0.88\text{MiB}$
- Total BRAM on board: 0.316MiB
- We need to store 4 images in memory!
- Let's use the ZBT RAM:  $2 \cdot 2.25\text{MiB}$
- Not so fast: 1 clock cycle per memory access



- 1 image is a lot of data:  
 $640 \cdot 480 \cdot 24 \text{ bits} \approx 0.88\text{MiB}$
- Total BRAM on board: 0.316MiB
- We need to store 4 images in memory!
- Let's use the ZBT RAM:  $2 \cdot 2.25\text{MiB}$
- Not so fast: 1 clock cycle per memory access
- 1 pixel per address would require a clock speed  $> 100\text{MHz}$

- 1 image is a lot of data:  
 $640 \cdot 480 \cdot 24 \text{ bits} \approx 0.88\text{MiB}$
- Total BRAM on board: 0.316MiB
- We need to store 4 images in memory!
- Let's use the ZBT RAM:  $2 \cdot 2.25\text{MiB}$
- Not so fast: 1 clock cycle per memory access
- 1 pixel per address would require a clock speed  $> 100\text{MHz}$
- Let's store store 18 bits per pixel or 2 per address

# memory\_interface: operation

1

# system io: ntsc\_capture

# system io: vga\_write

