

# Augmented Reality on FPGA

## Realtime Object Recognition and Image Processing

Logan Williams    José E. Cruz Serrallés

15 November 2011

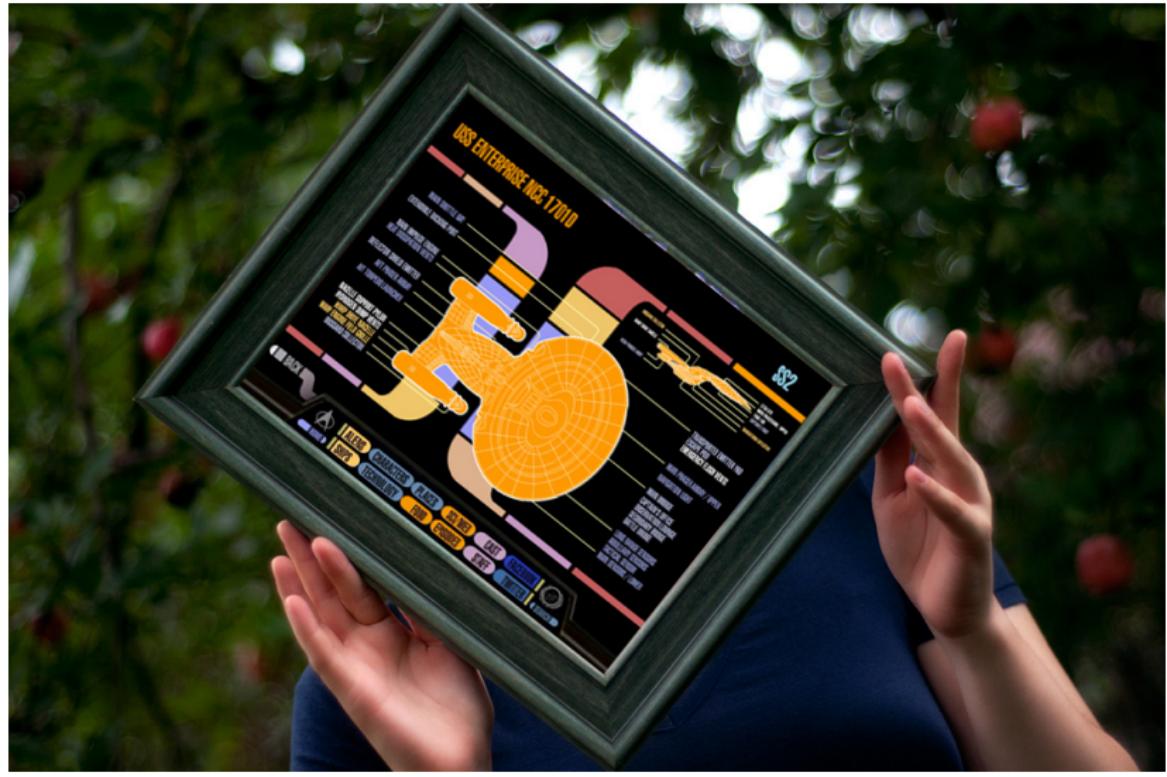
# Introduction

- Overlay a digital image on a physical object in realtime.
- In this case, we want to identify a picture frame in captured video, and output video with another image distorted to fit on top of the picture frame.

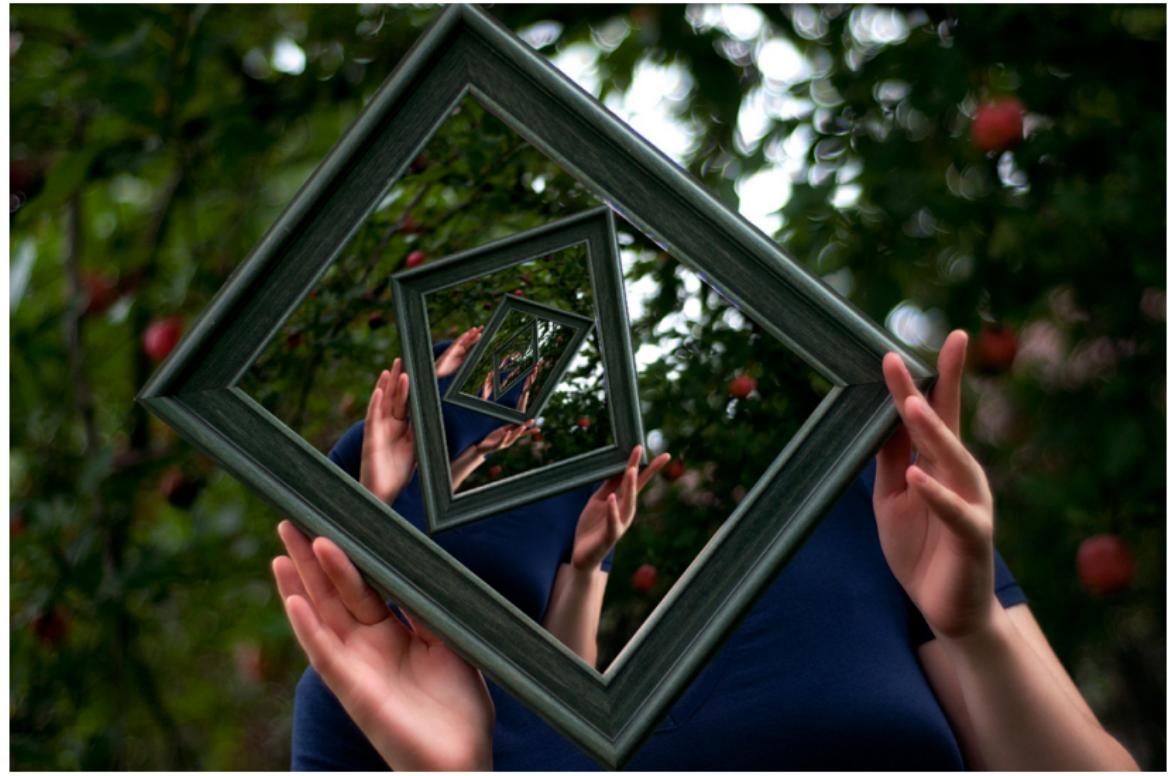
# Example Image



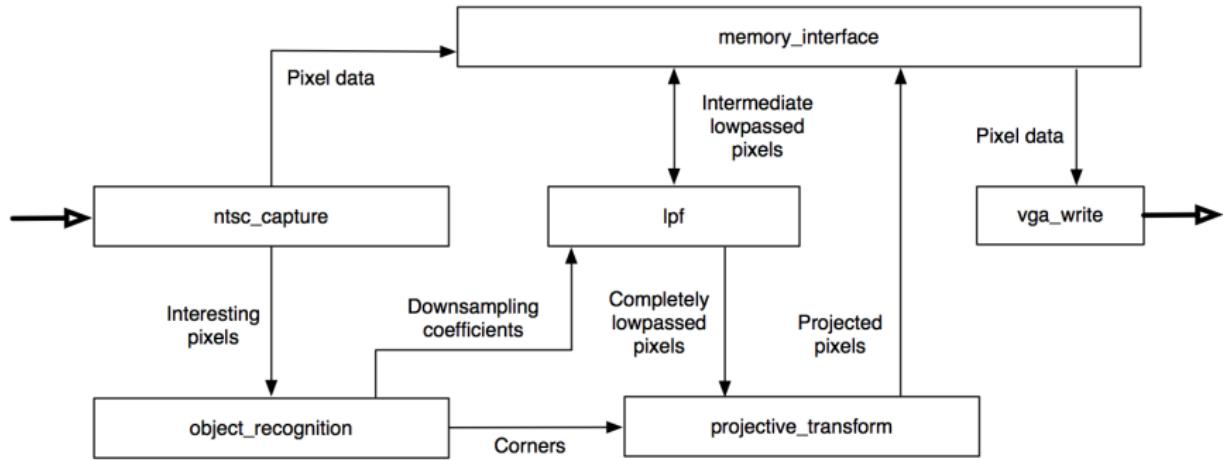
# Example Image



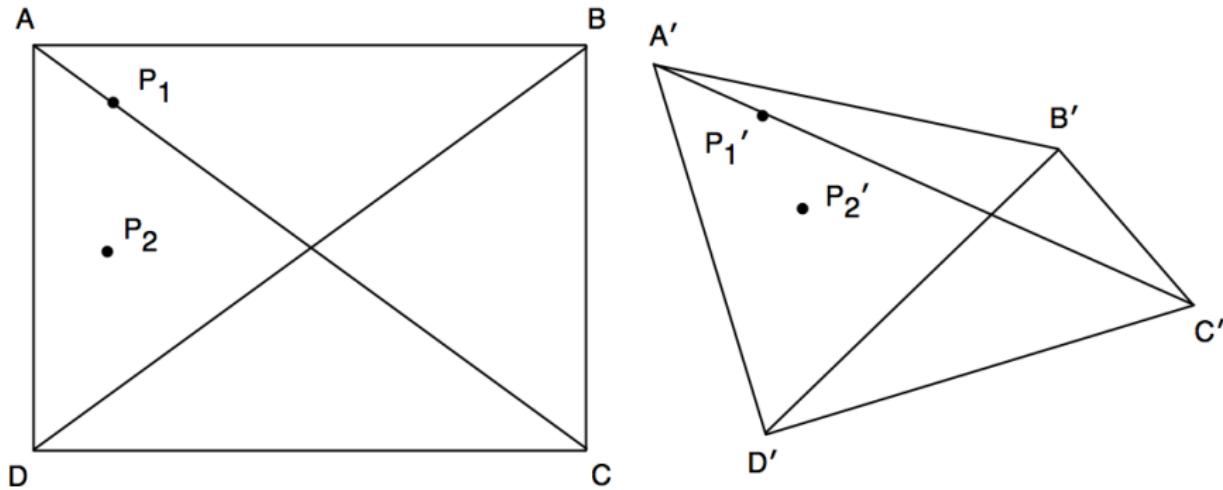
# Example Image



# Top-Level Overview

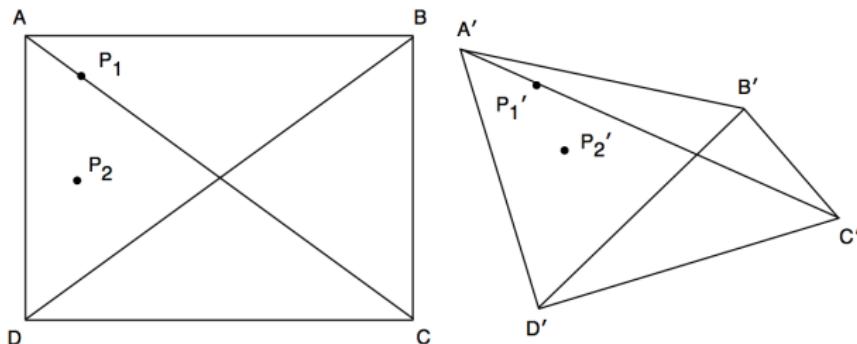


# projective\_transform: Purpose



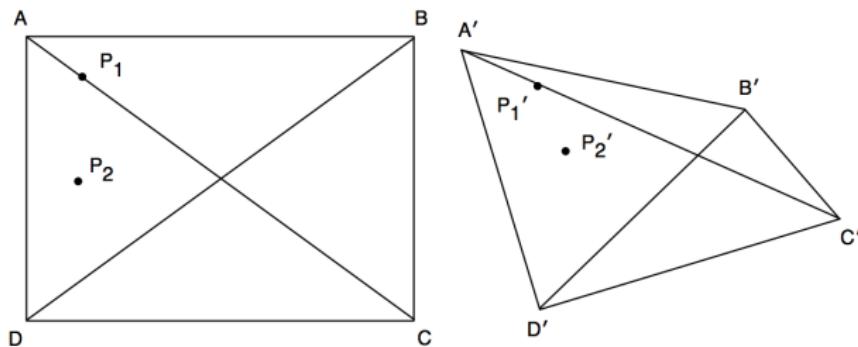
- Skew to any arbitrary convex quadrilateral

# projective\_transform: How the algorithm works



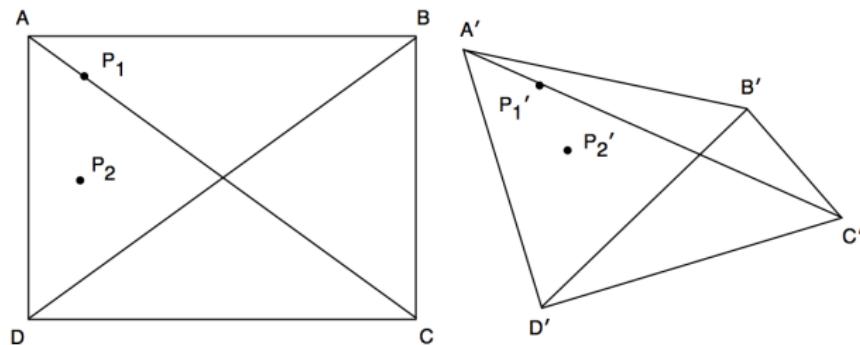
- 1 Calculate the distance of line  $\overline{A'D'}$  and assign it to  $d_{ad}$ .
- 2 Do the same for  $\overline{B'C'}$  and assign it to  $d_{bc}$ .
- 3 Create two “iterator points,” point  $I_A$  and  $I_B$  initially located at  $A'$  and  $B'$ .
- 4 Let  $o_x = 0$  and  $o_y = 0$
- 5 Calculate the distance between the iterator points, assign it to  $d_i$ .
- 6 Create a third iterator point,  $I_C$  at the location  $I_A$ .

# projective\_transform: How the algorithm works



- 7 Assign the pixel value of  $I_C$  to pixel  $(o_x, o_y)$  in the original image.
- 8 Move  $I_C$  along line  $\overline{I_A I_B}$  by an amount  $= \frac{d_i}{\text{width}_{\text{original}}}$ .
- 9 Increment  $o_x$ .
- 10 Repeat steps 7–9 until  $I_C = I_B$ .

# projective\_transform: How the algorithm works



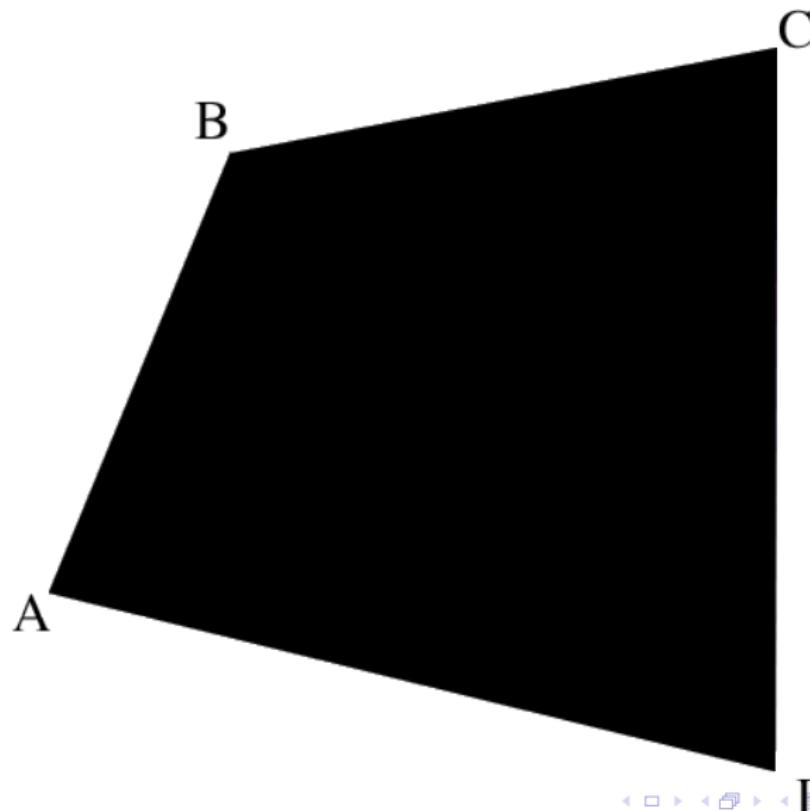
- 11 Move  $I_A$  along line  $\overline{A'D'}$  by an amount  $= \frac{d_{ad}}{\text{height}_{\text{original}}}$ .
- 12 Move  $I_B$  along line  $\overline{B'C'}$  by an amount  $= \frac{d_{bc}}{\text{height}_{\text{original}}}$ .
- 13 Increment  $o_y$ .
- 14 Repeat steps 5–13 until  $I_A = D'$  and  $I_B = C'$ .

# projective\_transform: Example

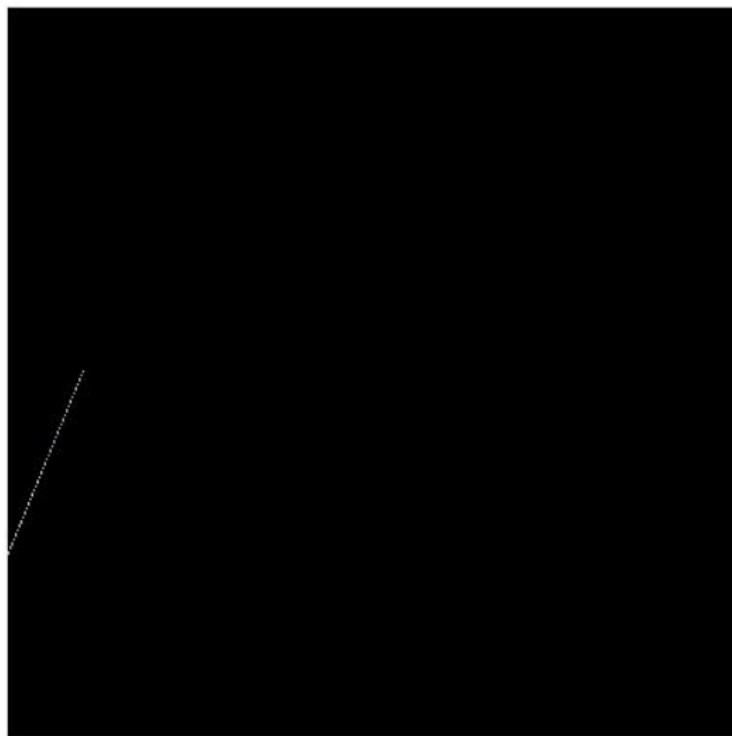


Figure: The original image

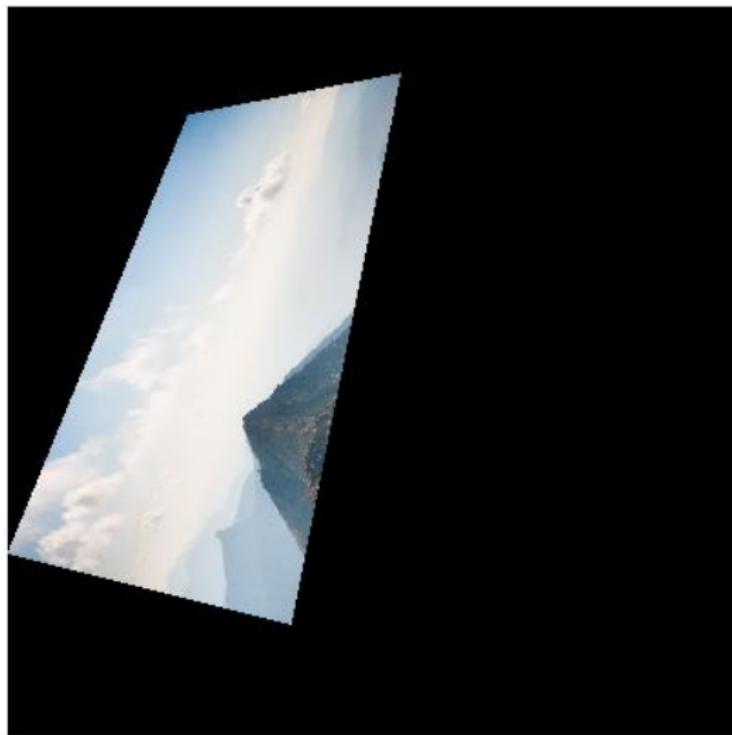
# projective\_transform: Example



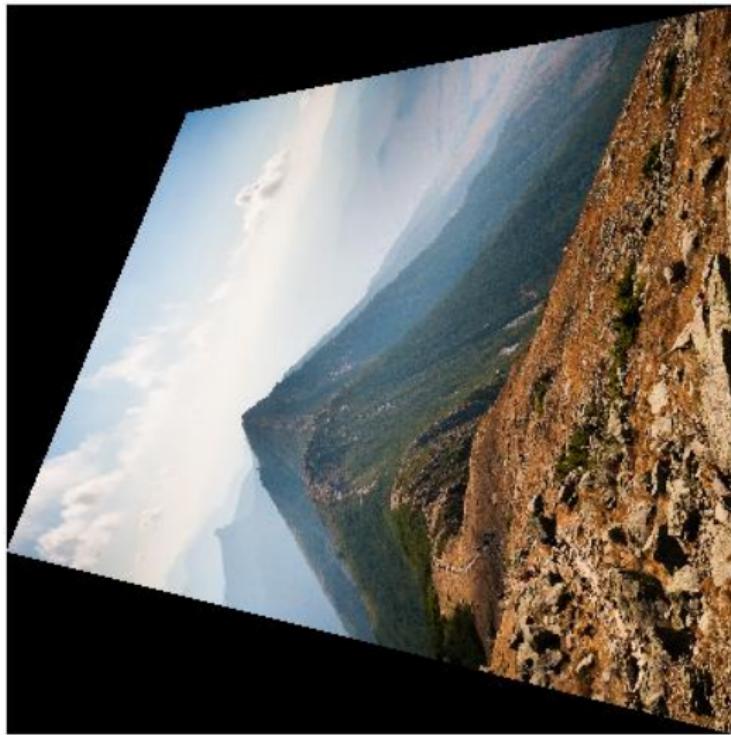
# projective\_transform: Example



# projective\_transform: Example



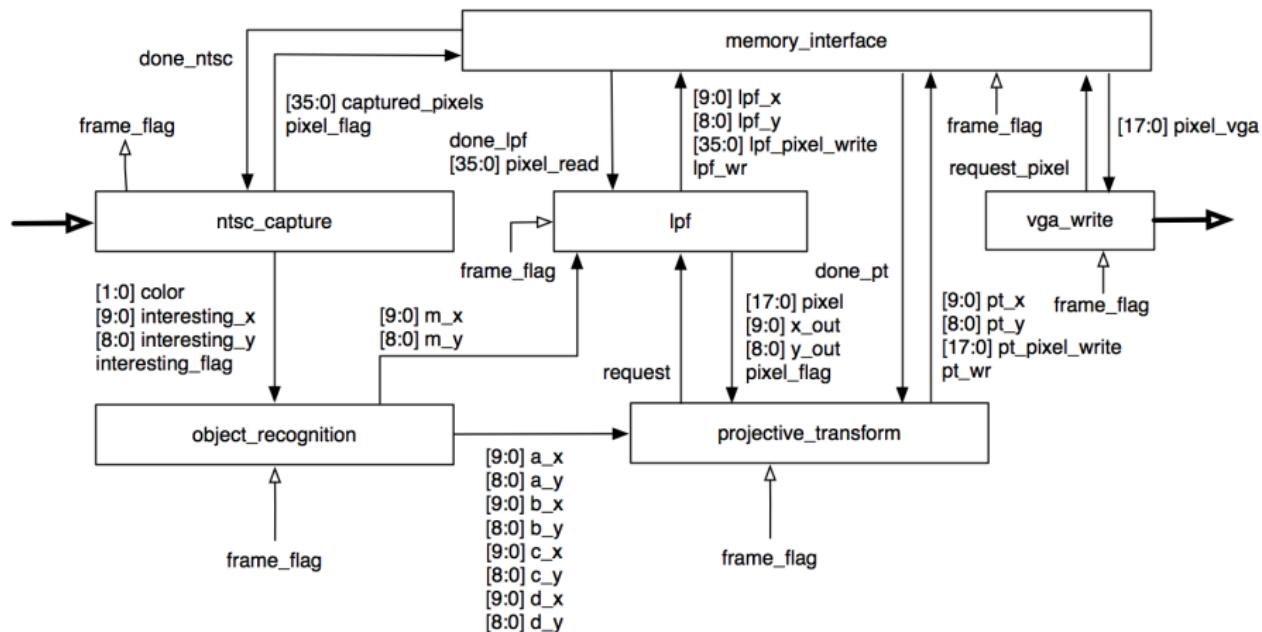
# projective\_transform: Example



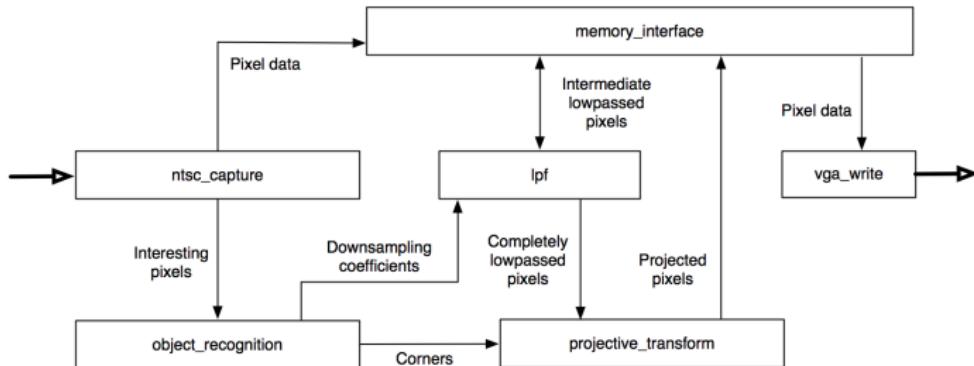
# projective\_transform: FPGA implementation

- Straightforward implementation of the above algorithm
- Uses coregen Divider modules for the divisions
- Requires only  $2*640*480 + 4*480$  multiplications per clock cycle
- Uses an iterative algorithm for finding distances (pipelined at the end of each line of the image)
- Processes pixels “on-the-fly” from LPF
- Negligible memory requirements (a handful of registers)

# projective\_transform: How it Interfaces



# object\_recognition



- Mark corners of frame with four differently colored dots.
- Recognition begins in the `ntsc_capture` module, which detects these colors as it is capturing data and sends the pixel info to the `object_recognition` module.

## object\_recognition

- Take linear weighted center of mass for each image
- Sums the (x,y) coordinates for each color as it receives them. (8 running sums, 2 for each color)
- When the frame is done, divide each sum by the number of summed items
- The resulting 4 (x,y) pairs are the corners of the frame
- By looking for pixels in ntsc\_capture we significantly reduce the amount of time spent in object\_recognition

# LPF: its purpose

- projective\_transform → aliasing



Original

# LPF: its purpose

- projective\_transform → aliasing
- Aliasing reduces the quality of an image

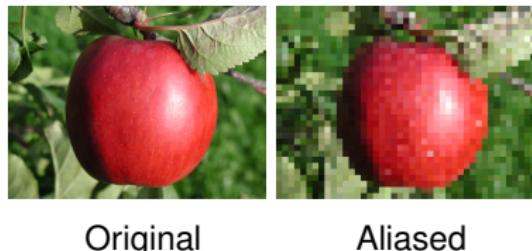


Original

Aliased

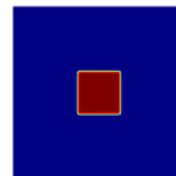
# LPF: its purpose

- `projective_transform` → aliasing
- Aliasing reduces the quality of an image
- Lowpass filtering prevents aliasing



Original

Aliased



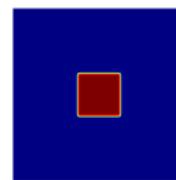
Mag. of Filter

# LPF: its purpose

- `projective_transform` → aliasing
- Aliasing reduces the quality of an image
- Lowpass filtering prevents aliasing



Original



Mag. of Filter



Filtered

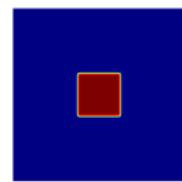
# LPF: its purpose

- `projective_transform` → aliasing
- Aliasing reduces the quality of an image
- Lowpass filtering prevents aliasing



Original

Aliased



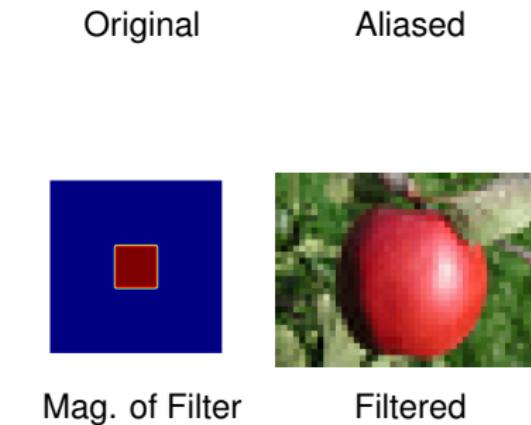
Mag. of Filter



Filtered

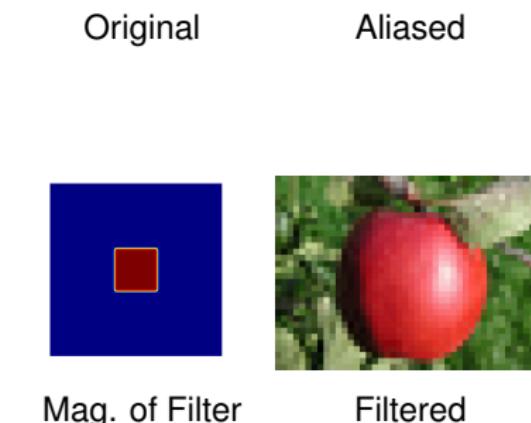
# LPF: its purpose

- `projective_transform` → aliasing
- Aliasing reduces the quality of an image
- Lowpass filtering prevents aliasing
- Information of an image is mostly phase



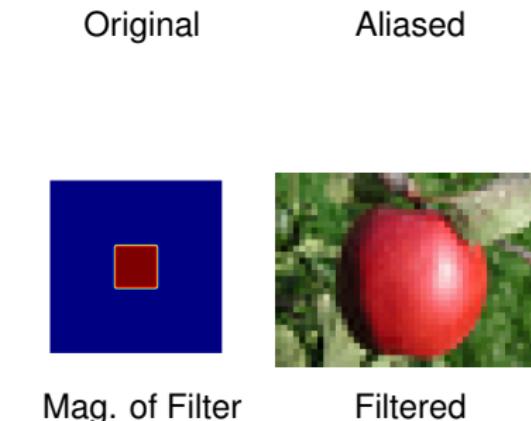
# LPF: its purpose

- `projective_transform` → aliasing
- Aliasing reduces the quality of an image
- Lowpass filtering prevents aliasing
- Information of an image is mostly phase
- Symmetric Type I FIR filter → 0 phase distortion



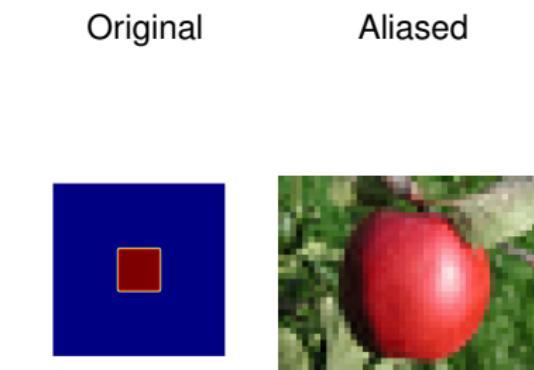
# LPF: its purpose

- `projective_transform` → aliasing
- Aliasing reduces the quality of an image
- Lowpass filtering prevents aliasing
- Information of an image is mostly phase
- Symmetric Type I FIR filter → 0 phase distortion
- Parks-McClellan: reasonable accuracy, symmetric, easily calculable



# LPF: its purpose

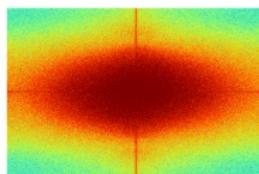
- projective\_transform → aliasing
- Aliasing reduces the quality of an image
- Lowpass filtering prevents aliasing
- Information of an image is mostly phase
- Symmetric Type I FIR filter → 0 phase distortion
- Parks-McClellan: reasonable accuracy, symmetric, easily calculable
- FIR PM filter reduces mem. acceses to 1.5/pixel



Mag. of Filter      Filtered

# LPF: the algorithm

- Given an arbitrary image & skewing coefficients  $M_x$  &  $M_y$ .



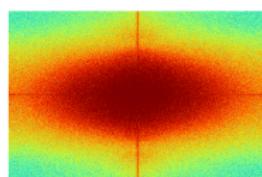
F.T. Original



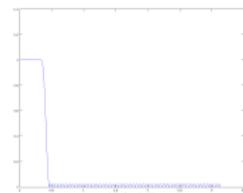
Original Image

# LPF: the algorithm

- ① Given an arbitrary image & skewing coefficients  $M_x$  &  $M_y$ .
- ② Fetch a filter with cutoff  $\frac{\pi}{M_y}$ .



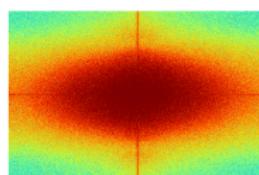
F.T. Original



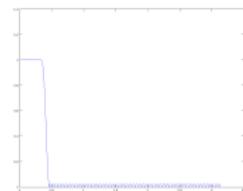
1D FIR,  $\omega_c = \frac{\pi}{8}$

# LPF: the algorithm

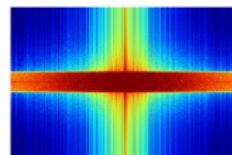
- ① Given an arbitrary image & skewing coefficients  $M_x$  &  $M_y$ .
- ② Fetch a filter with cutoff  $\frac{\pi}{M_y}$ .
- ③ Filter each column and store in memory.



F.T. Original



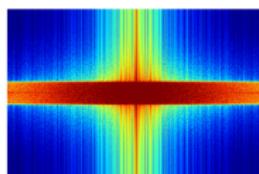
1D FIR,  $\omega_c = \frac{\pi}{8}$



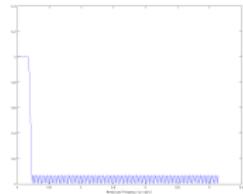
F.T. Filtered

# LPF: the algorithm

- ① Given an arbitrary image & skewing coefficients  $M_x$  &  $M_y$ .
- ② Fetch a filter with cutoff  $\frac{\pi}{M_y}$ .
- ③ Filter each column and store in memory.
- ④ Fetch a filter with cutoff  $\frac{\pi}{M_x}$ .



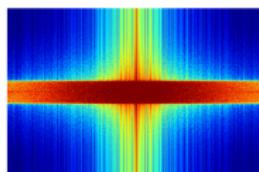
F.T. Filtered



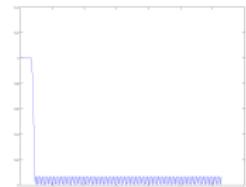
1D FIR,  $\omega_c = \frac{\pi}{16}$

# LPF: the algorithm

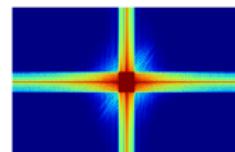
- ① Given an arbitrary image & skewing coefficients  $M_x$  &  $M_y$ .
- ② Fetch a filter with cutoff  $\frac{\pi}{M_y}$ .
- ③ Filter each column and store in memory.
- ④ Fetch a filter with cutoff  $\frac{\pi}{M_x}$ .
- ⑤ Filter each row and output to projective\_transform.



F.T. Filtered



1D FIR,  $\omega_c = \frac{\pi}{16}$



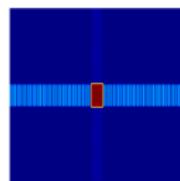
F.T. Output

# LPF: the algorithm

- 1 Given an arbitrary image & skewing coefficients  $M_x$  &  $M_y$ .
- 2 Fetch a filter with cutoff  $\frac{\pi}{M_y}$ .
- 3 Filter each column and store in memory.
- 4 Fetch a filter with cutoff  $\frac{\pi}{M_x}$ .
- 5 Filter each row and output to projective\_transform.
- 6 Repeat this process every refresh cycle.



Original



F.T. of Process



Output

# memory\_interface

- 1 image is a lot of data:  
 $640 \cdot 480 \cdot 24 \text{ bits} \approx 0.88\text{MiB}$

# memory\_interface

- 1 image is a lot of data:  
 $640 \cdot 480 \cdot 24 \text{ bits} \approx 0.88\text{MiB}$
- Total BRAM on board: 0.316MiB

# memory\_interface

- 1 image is a lot of data:  
 $640 \cdot 480 \cdot 24 \text{ bits} \approx 0.88\text{MiB}$
- Total BRAM on board: 0.316MiB
- We need to store 4 images in memory!

# memory\_interface

- 1 image is a lot of data:  
 $640 \cdot 480 \cdot 24 \text{ bits} \approx 0.88\text{MiB}$
- Total BRAM on board: 0.316MiB
- We need to store 4 images in memory!
- Let's use the ZBT RAM:  $2 \cdot 2.25\text{MiB}$

# memory\_interface

- 1 image is a lot of data:  
 $640 \cdot 480 \cdot 24 \text{ bits} \approx 0.88\text{MiB}$
- Total BRAM on board: 0.316MiB
- We need to store 4 images in memory!
- Let's use the ZBT RAM:  $2 \cdot 2.25\text{MiB}$
- Not so fast: 1 clock cycle per memory access

# memory\_interface

- 1 image is a lot of data:  
 $640 \cdot 480 \cdot 24 \text{ bits} \approx 0.88\text{MiB}$
- Total BRAM on board: 0.316MiB
- We need to store 4 images in memory!
- Let's use the ZBT RAM:  $2 \cdot 2.25\text{MiB}$
- Not so fast: 1 clock cycle per memory access
- 1 pixel per address would require a clock speed  $> 100\text{MHz}$

- 1 image is a lot of data:  
 $640 \cdot 480 \cdot 24 \text{ bits} \approx 0.88\text{MiB}$
- Total BRAM on board: 0.316MiB
- We need to store 4 images in memory!
- Let's use the ZBT RAM:  $2 \cdot 2.25\text{MiB}$
- Not so fast: 1 clock cycle per memory access
- 1 pixel per address would require a clock speed  $> 100\text{MHz}$
- Let's store 18 bits per pixel or 2 per address

# memory\_interface: operation

1

# system io: ntsc\_capture

# system io: vga\_write

# timeline